

Practical-1,2

Abhinav Sharma

Objective

The aim is to build some good linear models to analyze the **swiss** data where **Fertility** is used as the response.

Initial Data Analysis

Dataset Description

Standardized fertility measure and socio-economic indicators for each of 47 French-speaking provinces of Switzerland at about 1888.

Dataset Format

A data frame with 47 observations on 6 variables, each of which is in percent, i.e., in $[0, 100]$.

- [,1] **Fertility** - Ig, 'common standardized fertility measure'
- [,2] **Agriculture** - % of males involved in agriculture as occupation
- [,3] **Examination** - % draftees receiving highest mark on army examination
- [,4] **Education** - % education beyond primary school for draftees.
- [,5] **Catholic** - % 'catholic' (as opposed to 'protestant').
- [,6] **Infant.Mortality** - live births who live less than 1 year.

All variables but 'Fertility' give proportions of the population.

Analysis


```
# loading data
data(swiss)

# summary of data
summary(swiss)
```

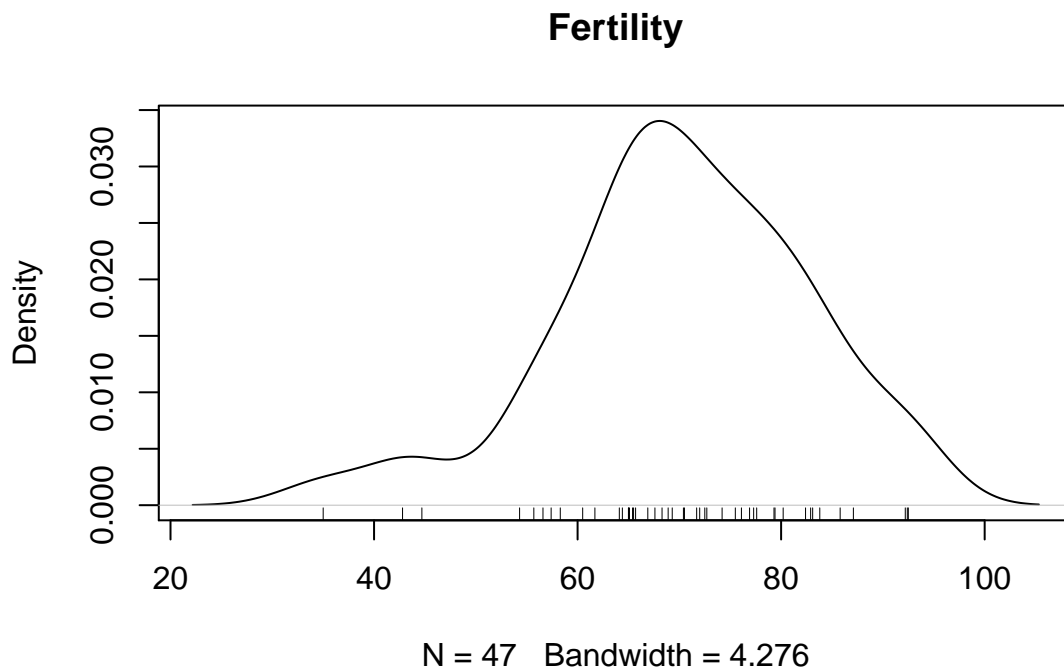
```
##      Fertility      Agriculture      Examination      Education
##  Min.   :35.00    Min.    : 1.20    Min.     : 3.00    Min.     : 1.00
## 1st Qu.:64.70    1st Qu.:35.90    1st Qu.:12.00    1st Qu.: 6.00
## Median :70.40    Median :54.10    Median :16.00    Median : 8.00
## Mean   :70.14    Mean   :50.66    Mean   :16.49    Mean   :10.98
## 3rd Qu.:78.45    3rd Qu.:67.65    3rd Qu.:22.00    3rd Qu.:12.00
## Max.   :92.50    Max.    :89.70    Max.    :37.00    Max.    :53.00
##      Catholic      Infant.Mortality
##  Min.    : 2.150    Min.     :10.80
## 1st Qu.: 5.195    1st Qu.:18.15
## Median :15.140    Median :20.00
## Mean    :41.144    Mean     :19.94
## 3rd Qu.:93.125    3rd Qu.:21.70
## Max.    :100.000    Max.     :26.60
```

1. Distribution of response variable

```
# histogram plot  
hist(swiss$Fertility,main="Fertility",xlab="Fertility")
```




```
# density plot  
plot(density(swiss$Fertility),main="Fertility")  
rug(swiss$Fertility)
```



- Plot is little right skewed
- Fertility rates are mostly between 60-90%

2. Analyzing predictor variables

```
# boxplot analysis  
boxplot(swiss)
```



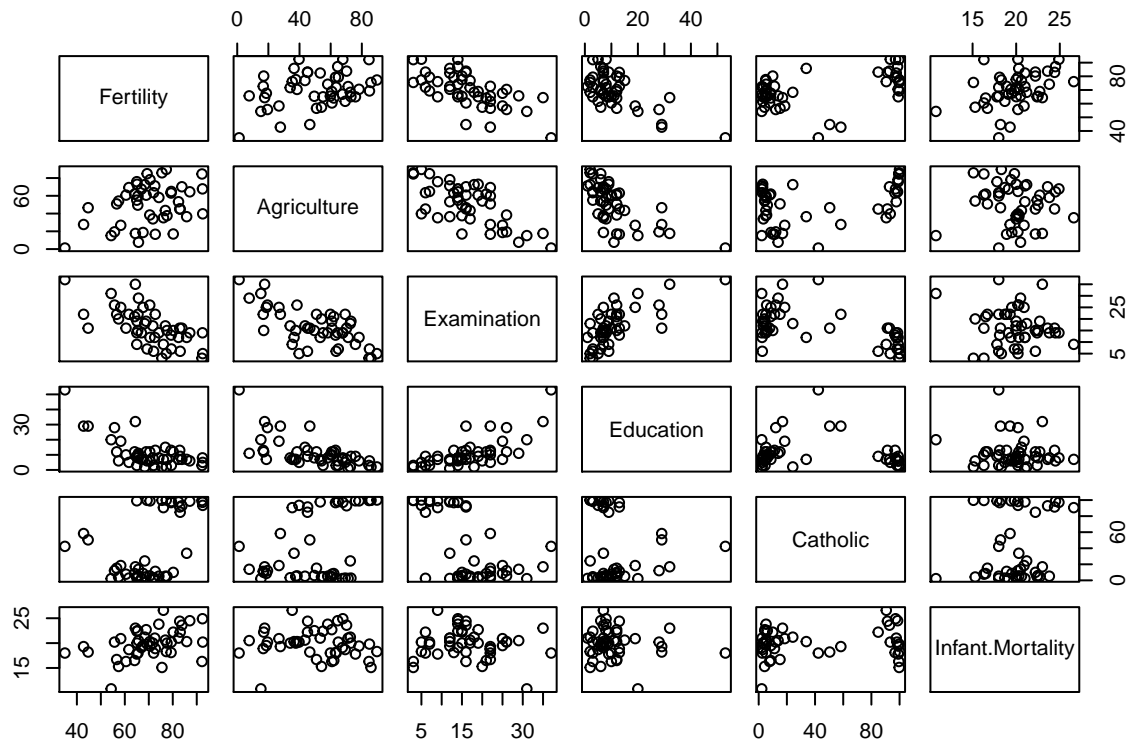
- Catholic variable covers wide range of values
- Infant.Mortality variable is very condensed
- Education seems to have some outliers

```
# correlation analysis  
cor(swiss)
```

```
##           Fertility Agriculture Examination Education Catholic  
## Fertility    1.0000000  0.35307918 -0.6458827 -0.66378886  0.4636847  
## Agriculture  0.3530792  1.00000000 -0.6865422 -0.63952252  0.4010951  
## Examination -0.6458827 -0.68654221  1.0000000  0.69841530 -0.5727418  
## Education   -0.6637889 -0.63952252  0.6984153  1.00000000 -0.1538589  
## Catholic    0.4636847  0.40109505 -0.5727418 -0.15385892  1.0000000  
## Infant.Mortality 0.4165560 -0.06085861 -0.1140216 -0.09932185  0.1754959  
##           Infant.Mortality  
## Fertility      0.41655603  
## Agriculture    -0.06085861  
## Examination    -0.11402160  
## Education      -0.09932185  
## Catholic       0.17549591  
## Infant.Mortality 1.00000000
```

- All correlations with Fertility are less than 0.8, indicating no signs of strong multicollinearity.
- Correlations are between 0.3-0.8, indicating mild multicollinearity.

```
# correlation analysis plot
pairs(swiss)
```



- Plot shows linear relationship between Agriculture and Examination.
- Also, between Examination and Education.
- Interpretation of coefficients will be affected.

Fitting a linear Model

```
# basic linear model with all variables
model <- lm(Fertility ~ Agriculture + Examination + Education +
            Catholic + Infant.Mortality, swiss)
summary(model)
```


```
##
## Call:
## lm(formula = Fertility ~ Agriculture + Examination + Education +
##     Catholic + Infant.Mortality, data = swiss)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -15.2743  -5.2617   0.5032   4.1198  15.3213
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   66.91518   10.70604    6.250 1.91e-07 ***
## Agriculture   -0.17211    0.07030   -2.448  0.01873 *
## Examination   -0.25801    0.25388   -1.016  0.31546
## Education     -0.87094    0.18303   -4.758 2.43e-05 ***
## Catholic       0.10412    0.03526    2.953  0.00519 **
```

```
## Infant.Mortality 1.07705    0.38172    2.822  0.00734 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 7.165 on 41 degrees of freedom
## Multiple R-squared:  0.7067, Adjusted R-squared:  0.671
## F-statistic: 19.76 on 5 and 41 DF,  p-value: 5.594e-10
```

Transformations

We can use box-cox transformation method to understand if response variable transformation is needed or not.

```
# box-cox transformation for response variable
library(MASS)
box <- boxcox(
  Fertility ~ Agriculture + Examination + Education + Catholic +
  Infant.Mortality, data = swiss,
  lambda = seq(from = -3, to = 3, length = 50)
)
```



- $\lambda=1$ lies in 95% confidence interval, so we can say that no transformation is needed for response variable.

Transformations of predictors can also be achieved - ?? Maybe polynomials ??

Variable Selection

```
# stepwise selection using AIC
model2 <- step(model1, trace=FALSE)
model2

##
## Call:
## lm(formula = Fertility ~ Agriculture + Education + Catholic +
```

```
## Infant.Mortality, data = swiss)
##
## Coefficients:
## (Intercept)      Agriculture      Education      Catholic
##      62.1013      -0.1546      -0.9803      0.1247
## Infant.Mortality
##      1.0784
```

- The final AIC of model achieved = 189.86
- Final Model: Fertility ~ Agriculture + Education + Catholic + Infant.Mortality

```
# backward elimination using F-test
```

```
drop1(model,test='F')
```

```
## Single term deletions
```

```
##
```

```
## Model:
```

```
## Fertility ~ Agriculture + Examination + Education + Catholic +
```

```
## Infant.Mortality
```

	Df	Sum of Sq	RSS	AIC	F value	Pr(>F)
<none>			2105.0	190.69		
Agriculture	1	307.72	2412.8	195.10	5.9934	0.018727 *
Examination	1	53.03	2158.1	189.86	1.0328	0.315462
Education	1	1162.56	3267.6	209.36	22.6432	2.431e-05 ***
Catholic	1	447.71	2552.8	197.75	8.7200	0.005190 **
Infant.Mortality	1	408.75	2513.8	197.03	7.9612	0.007336 **

```
## ---
```

```
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

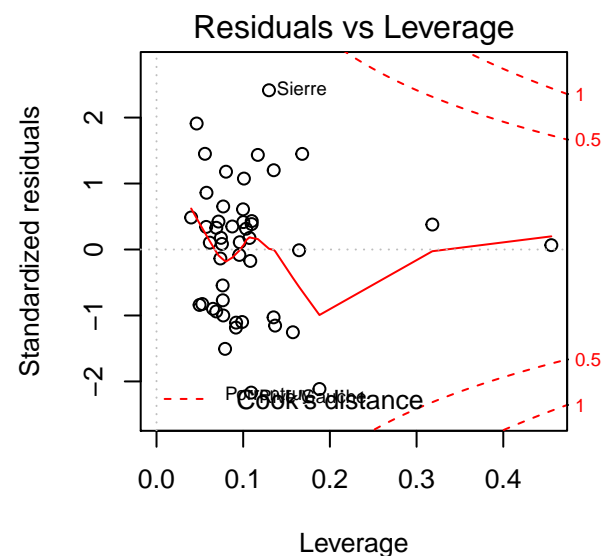



- Backward elimination using F-test also suggest removal of Examination term as p-value is extremely large (0.315462 > 0.05).
- Final Model is same as selected using stepwise.

Diagnostics

```
# Basic diagnostic plots
```

```
par(mfrow=c(2,2))
```

```
plot(model2)
```




- From residual plot, there is no pattern observed so constant variance assumption holds.
- From QQ plot, we don't see heavy tails so normality assumption also holds.
- Some extreme values are observed in cook's distance plot, indicating outliers and influential points in the data.

top influential points

```
swiss[cooks.distance(model2) > 0.1,]
```

```
##           Fertility Agriculture Examination Education Catholic
## Porrentruy      76.1         35.3           9           7     90.57
## Sierre         92.2         84.6           3           3     99.46
## Rive Gauche    42.8         27.7          22          29     58.33
##
##           Infant.Mortality
## Porrentruy             26.6
## Sierre                16.3
```

```
## Rive Gauche          19.3
# possible outliers
library(faraway)
halfnorm(influence(model2)$hat)
```



```
swiss[influence(model2)$hat>0.3,]
```

```
##           Fertility Agriculture Examination Education Catholic
## La Vallee      54.3         15.2           31          20      2.15
## V. De Geneve   35.0          1.2           37          53     42.34
##
##           Infant.Mortality
## La Vallee           10.8
## V. De Geneve        18.0
```

- We can handle this by either removing outliers points from the dataset and then proceed with least squares.
- We can also perform robust regression that downweights the effects of larger errors.

```
# Robust regression
library(MASS)
rlmodel <- rlm(Fertility ~ Agriculture + Examination + Education +
               Catholic + Infant.Mortality, swiss)
summary(rlmodel)
```

```
##
## Call: rlm(formula = Fertility ~ Agriculture + Examination + Education +
##           Catholic + Infant.Mortality, data = swiss)
## Residuals:
##      Min       1Q   Median       3Q      Max
## -16.0173  -5.0005   0.3231   4.0905  16.1854
##
## Coefficients:
##              Value      Std. Error t value
```



```
## (Intercept)      65.4851 10.9074      6.0037
## Agriculture      -0.1933  0.0716     -2.6994
## Examination      -0.2869  0.2587     -1.1091
## Education        -0.8495  0.1865     -4.5558
## Catholic          0.1049  0.0359      2.9209
## Infant.Mortality  1.2184  0.3889      3.1330
##
## Residual standard error: 6.633 on 41 degrees of freedom
```

Predictions

Suppose we assume following data for a particular region:

- [,2] Agriculture - 70 % of males involved in agriculture as occupation
- [,3] Examination - 20 % draftees receiving highest mark on army examination
- [,4] Education - 30 % education beyond primary school for draftees.
- [,5] Catholic - 10 % 'catholic' (as opposed to 'protestant').
- [,6] Infant.Mortality - 28 live births who live less than 1 year

```
# input data
data <- data.frame(Agriculture=70,Examination=20,Education=30,
                   Catholic=10,Infant.Mortality=28)
# predict fertility rate for a region using full model
predict(model, newdata=data)
```

```
##          1
## 54.77734
```

```
# predict fertility rate using stepwise model
predict(model2, newdata=data)
```

```
##          1
## 53.31322
```


```
# predict fertility rate using robust regression model
predict(rlmodel, newdata=data)
```

```
##          1
## 55.89241
```

Interpretation

Let's use summary statistics of different variables to get different values of predictions and then interpret results for fertility.

```
# input data
data <- data.frame(
  Agriculture=c(1.20,35.90,50.66,67.95,89.70),
  Examination=c(3,12,16.49,22,37),
  Education=c(1,6,10.98,12,53),
  Catholic=c(2.150,5.195,41.144,93.125,100),
  Infant.Mortality=c(10.80,18.15,19.94,21.70,26.60)
)
# using full model to get complete interpretation
data$pp <- predict(model,newdata=data)
matplot(data,type='o',ylab='Fertility',xlab='Variables')
```



Here:

- 1-Agriculture
- 2-Examination
- 3-Education
- 4-Catholic
- 5-Infant.Mortality
- 6-Fertility

Interpretation:

- Fertility goes down as other predictor variables % increases
- Fertility is highly affected by **Agriculture** and **Catholic** variables.